



Introduction to Operating Systems

Prepared by:
Zeyad Ashraf



I/O Management

Content

- Need of I/O Management
- I/O Subsystem
- I/O Devices Categories:
 - Block Devices
 - Character Devices
- I/O Protocols Categories:
 - Special Instructions I/O
 - Memory-Mapped I/O
 - Direct Memory Access (DMA)
- Interrupt Handling Mechanisms
- Synchronous Vs. Asynchronous I/O
- Synchronization and Critical Sections
 - Mutex
 - Semaphore
- Deadlocks


Course Outlines 

Course Outlines

4.1 Need for I/O Management

- As part of the system, there could be multiple devices that are connected and perform different **input-output functions**.
- These I/O devices could be used for **human interaction** such as **display panel, touch panels, keyboard, mouse, and track pads**, to name a few.
- Another I/O devices could be to **connect the system to storage devices, sensors**, and so on.
- There could also be I/O devices for **networking** needs that implement certain parts of the networking stack. These could be **Wi-Fi, Ethernet, and Bluetooth** devices and so on.
- They vary from one to another in the form of **protocols** they use to communicate such as the **data format, speed** at which they operate, error reporting mechanisms, and more.
- The OS presents a unified I/O system that **abstracts** the complexity from applications. The OS handles this by establishing protocols and interfaces with each I/O controller. However, the I/O subsystem usually forms the complex part of the operating system due to the dynamics and the wide variety of I/Os involved.

4.1 Need for I/O Management



4.2 I/O Subsystem

- Input/output devices that are connected to the computer are called **peripheral** devices.
- It is communicated with the system through the busses:
 - **Data bus**: to transfer data
 - **Address bus**: used to specify address locations, and
 - **Control bus**: to control a device.
- There **could be different buses or device protocols** that an OS may support.
- The most common protocols include:
 - Peripheral Component Interconnect Express (PCIe) protocol,
 - Inter-Integrated Circuit (I2C), and
 - Advanced Configuration and Power Interface (ACPI)
- **A device** can be connected over **one or more** of these **interfaces**.

4.2 I/O Subsystem

- Typically, there is a software component in kernel mode called as the “**device driver**” that handles all interfaces with a device.
- It helps with communicating between the device and the OS and **abstracts** the device specifics.
- Similarly, there could be a driver at the bus level usually referred to as **the bus driver**. (Ex: USB Bus driver)
- Most OSs include an inbox driver that implements the bus driver.
- There is usually a driver for each controller and each device.

4.3 I/O Devices Categories

- The I/O devices can be broadly divided into two categories called **block** and **character** devices.

1. **Character:** data represented sequentially as characters **ex..** Keyboard ,mouse
2. **Block:** data represented as blocks **ex..** CD, hard disk and flash.

- The **driver** would fill the required data and issue a command.
- The device **firmware** would respond back to the command and return a code that is utilized by the driver.
- The protocol, size, and format could differ from one device to another.

4.3.1 Block Devices



- These are devices with which the I/O device controller communicates by sending blocks of data.
- A block is referred to as a group of bytes that are referred together for Read/Write purposes.
- Example: **flash memory, digital camera**
- The device driver would access by specifying the size of Read/Writes which is **varying** from device to another.

4.3.2 Character Devices



- Another class of devices are character devices, the subtle difference is that the communication happens by sending and receiving single characters, which is usually a byte or an octet.
- Many serial port devices like **keyboards**, some **sensor devices**, and microcontrollers follow this mechanism.

4.4 I/O Protocols Categories




- The protocols used by the different devices (block devices or character devices) **could vary from one to another.**
- There are three main categories of I/O protocols that are used:
 1. Special Instruction I/O
 2. Memory-Mapped I/O
 3. Direct Memory Access (DMA)

4.4.1 Special Instruction I/O

- There could be specific **CPU instructions** that are custom developed for communicating with and controlling the I/O devices.
- Each device has a **unique I/O address**
- For example, there could be a CPU-specific protocol to communicate with the **embedded controller**.
- This may be needed for **faster and efficient** communication.
- However, such type of I/Os are special and smaller in number.


4.4.1 Special Instruction I/O




4.4.2 Memory-Mapped I/O

- The **most common** form of I/O protocol is memory-mapped I/O (MMIO).
- The device and OS agree on **a common address range** carved out by the OS, and the **I/O device** makes reads and writes from/to this space to communicate to the OS.
- OS components such as drivers will communicate using this **interface** to talk to the device.
- Hence, it is used to enable **high-speed communication** for **network and graphics devices** that require high data transfer rates due to the volume of data being passed.

4.4.2 Memory-Mapped I/O



4.4.2 Memory-Mapped I/O



4.4.3 Direct Memory Access (DMA)

- There could be **devices that run at a slower speed than CPU** or the bus it is connected on. In this case, the device can leverage DMA.
- Here, the **OS grants authority to another controller**, usually referred to as the direct memory access controller, to **interrupt** the CPU after a specific data transfer is complete.
- The devices running at a smaller rate **can communicate back to the DMA controller** after completing its operation.
 - As a programmer, you could be interacting with devices that may perform **caching** and have different **error reporting** mechanisms, protocols, and so on.
- **Ex :** Video & audio devices, hard disks .

4.5 Interrupt Handling Mechanisms



- Polling or programmed I/O:
 - In this mechanism, **each period of time examine interrupt** information and calls a **specific routine** (driver)
- Interrupt-Driven I/O:
 - In This mechanism, there is an **interrupt vector** which contains the addresses for all **Interrupt Service Routines** (ISR) –**drivers**- for all I/O devices, according the **Interrupt Request Number** (IRQ) –which **unique** for each device- the OS acquire the address of the address of the correct service routine
 - Most of OSs using this mechanism

4.5 Interrupt Handling Mechanisms




- **Synchronous I/O: Example:** Printing Operation
 - Process request I/O operation
 - I/O operation is started
 - I/O Operation is complete
 - Control is returned to the user process
- **Asynchronous I/O:** Example: Using Networking sending and receiving using threads
 - Process Request I/O operation
 - I/O operation is started
 - Control is returned immediately to the user process
 - I/O continues while system operations occur

4.7 Synchronization and Critical Sections



- In multi-threaded applications, if one thread tries to change the value of **shared data** at the **same time** as another thread tries to read the value, there could be a **race condition** across threads.
- In this case, the result can be **unpredictable**.
- The access to such shared variables via shared memory, files, ports, and other I/O resources **needs to be synchronized** to protect it from being corrupted.
- In order to support this, the operating system provides **mutexes** and **semaphores** to coordinate access to these shared resources.


4.7 Synchronization and Critical Sections



4.7.1 Mutex




- A mutex is used for implementing ***mutual exclusion***: either of the participating processes or threads can have the key (mutex) and proceed with their work.
- **The other one would have to wait until the one holding the mutex finishes**



4.7.2 Semaphore




- A semaphore is a **generalized mutex**.
- A binary semaphore can assume a value of 0/1 and can be used to perform locks to certain critical sections.



4.8 Deadlocks

- When a set of processes become blocked because each process is holding a resource and waiting for another resource acquired by some other process. This is called as a **deadlock**.
- Process A holds Resource 1 and requires Resource 2. However, Process B already is holding Resource 2, but requires Resource 1. Unless either of them releases their resource.
- either of the processes may be able to move forward with the execution.**



4.8 Deadlocks



- A deadlock can arise if the following four conditions hold:
 - **Mutual Exclusion:** There is at least one resource on the system that is **not shareable**. This means that only one process can access this at any point in time. In the preceding example, Resources 1 and 2 can be accessed by only one process at any time.
 - **Hold and Wait:** A process is **holding at least one resource** and is waiting for other resources to proceed with its action. In the preceding example, both Processes A and B are holding at least one resource.
 - **No Preemption:** A resource cannot be **forcefully** taken from a process unless released automatically.
 - **Circular Wait:** A set of processes are **waiting for each other in circular form**.

Thank You

With My Best Wishes

Zeyad ashraf